

Node-screening tests for the ℓ_0 -penalized least-squares problem

Théo Guyard*, Cédric Herzet†, Clément Elvira‡

*Applied Mathematics Department, INSA Rennes, France | †SIMSMART team, INRIA Rennes-Bretagne Atlantique, France | ‡SCEE team, CentraleSupélec Rennes, France

Objectives

Reduce the optimization time of a Branch-and-Bound (BnB) algorithm tailored to the ℓ_0 -penalized least-squares problem by detecting nodes of the decision tree that cannot yield a global optimizer.

Introduction

- **Sparse decomposition:** Find an approximation of a vector \mathbf{y} as the *linear combination* of a *few* columns of a matrix \mathbf{A} .

ℓ_0 -penalized least-squares

$$\begin{array}{cc} \|\mathbf{y} - \mathbf{A}\mathbf{x}\|_2 & \|\mathbf{x}\|_0 \\ \text{Ensures data-fidelity} & \text{Promotes sparsity} \end{array}$$

$$(\mathcal{P}) \quad \begin{cases} \min & \frac{1}{2}\|\mathbf{y} - \mathbf{A}\mathbf{x}\|_2^2 + \lambda\|\mathbf{x}\|_0 \\ \text{st.} & \|\mathbf{x}\|_\infty \leq M \end{cases}$$

- **Characteristics:**
 - NP-hard problem due to the ℓ_0 -norm
 - Mixed-Integer Program reformulation
 - Addressable with BnB algorithms
 - “Big-M” to construct bounded relaxations
- **Recent advances by Atamturk *et. al.*:**
 - **Screening tests** to detect zeros and non-zeros in the optimizers of the problem
 - Dimensionality reduction in **pre-processing**
- **Our contributions:**
 - **Node-screening tests** to detect zeros and non-zeros in the optimizers of **any node** problem
 - Dimensionality reduction at **any step** of the optimization process
 - Larger improvement in the solving time

BnB algorithms

- **Generic procedure:**

Algorithm 1: BnB algorithm


Create a root node ν^0 and initialize $\mathcal{Q} = \{\nu^0\}$
while $\mathcal{Q} \neq \emptyset$ **do**
 1) Select a new node ν in \mathcal{Q}
 2) Find a lower and an upper bound on the node problem objective value
 3) Update the best upper bound known
 if $\text{node } LB \leq \text{best } UB$ **then**
 4) Create child nodes of ν by fixing new constraints and push them in \mathcal{Q}
 end
 5) Remove ν from \mathcal{Q}
end
return any node yielding the best UB

- **Efficiency of the algorithm:**

- Tightness of the bounds computed
- Ability to process nodes quickly
- **Number of nodes processed**

- **Particularized to our problem:**


- Binary tree with decision about the **nullity of an entry** at each node
- Each node is uniquely defined by the set of entries forced to **zero** and to **non-zero**



Node-screening tests

- **Underlying idea:**

- For many nodes, it is obvious that they cannot yield a global solution (too large penalty, *etc.*)
- BnB is an exact method so we have to **prove** that they can be pruned safely
- Standard pruning methodology : solve a convex relaxation. This is **expensive** !
- Why not trying **weaker** but **cheaper** methods ?

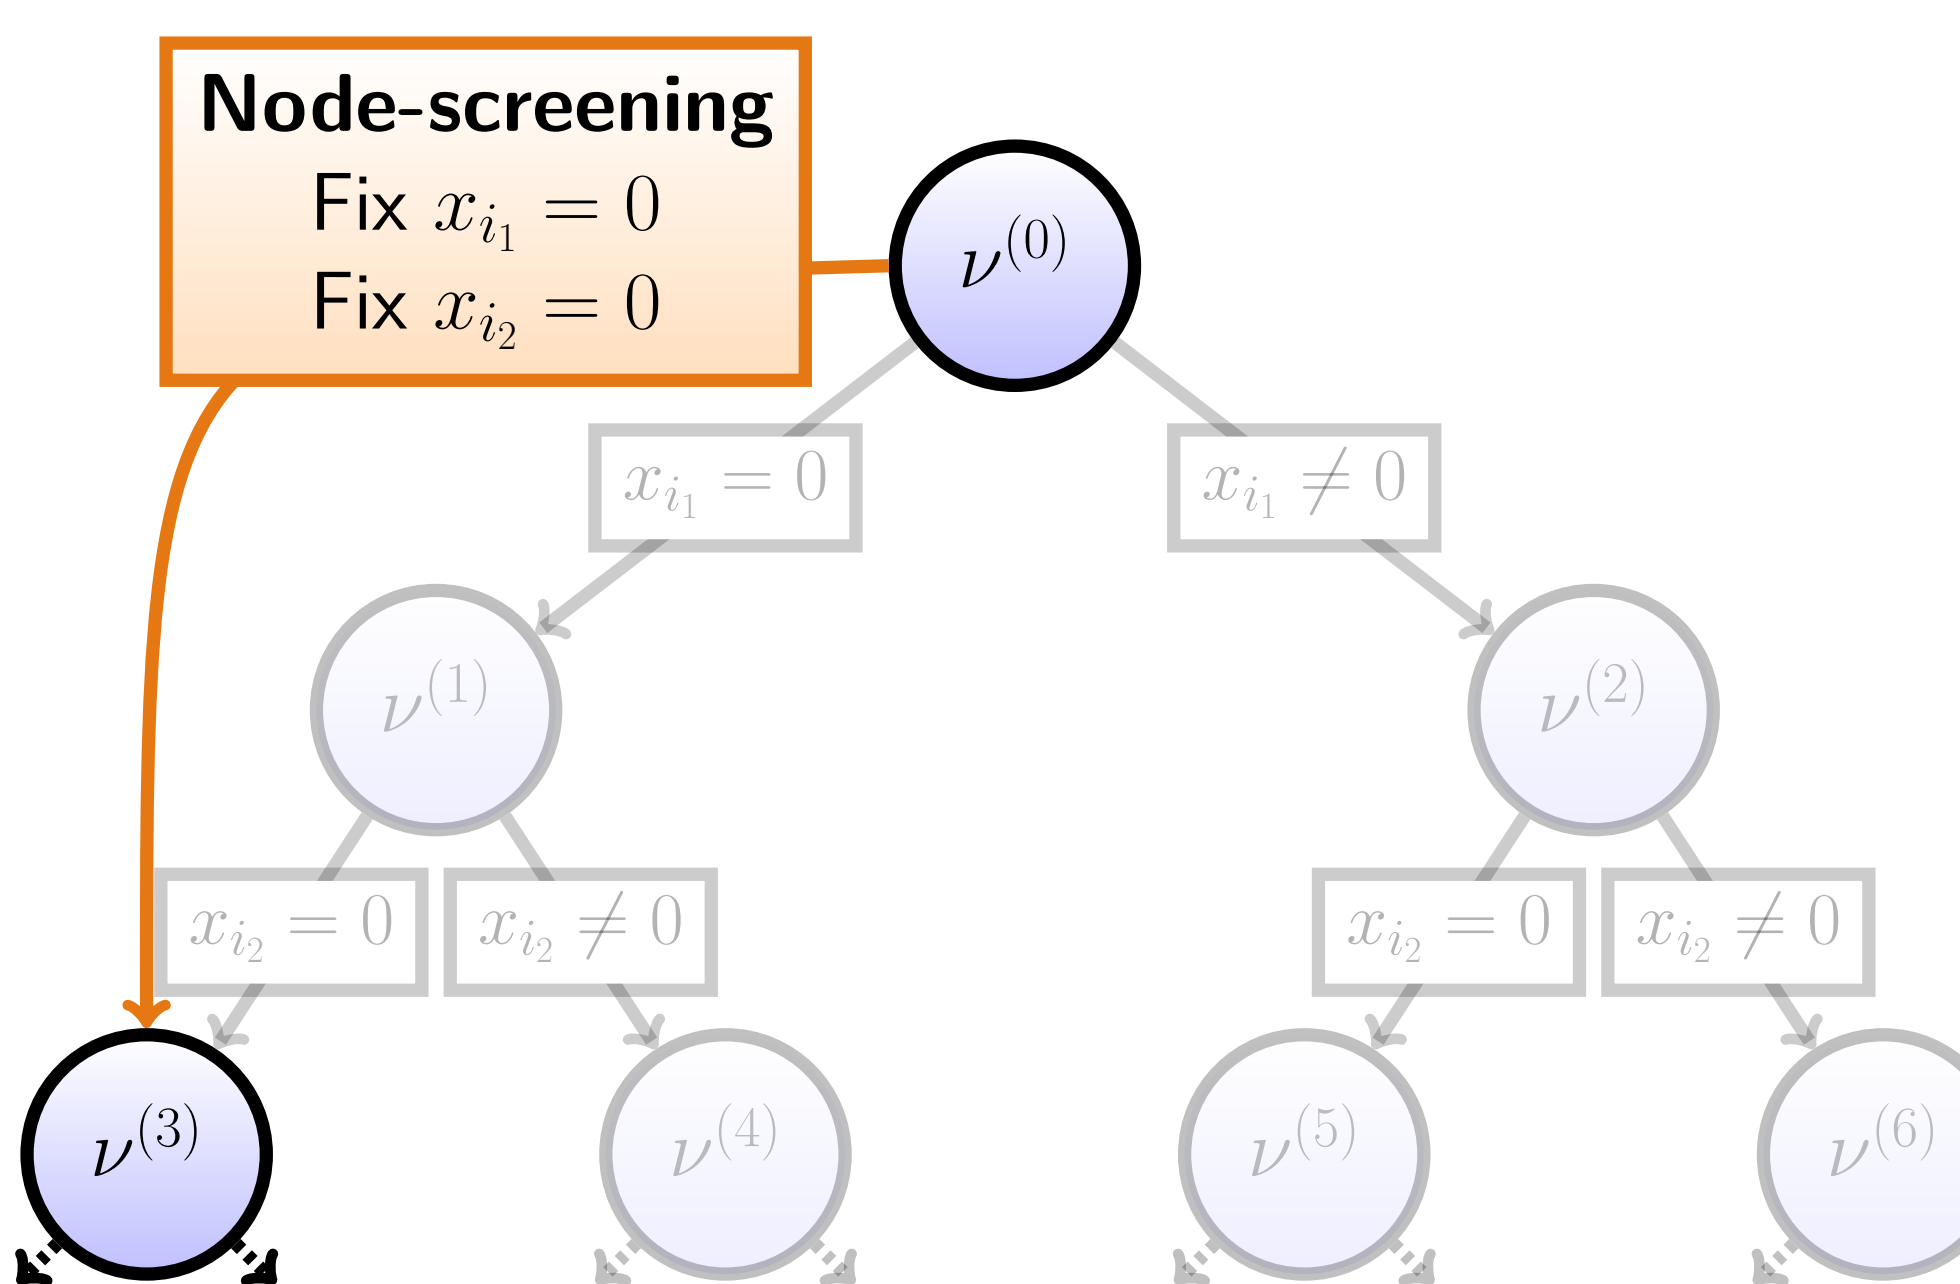


- **Differences with the classical method:**

- Test **any** unfixed index, no branching rule
- No computational overhead to prune nodes
- Can fix **multiple variables simultaneously**

Main ingredients

- At each node, a **convex relaxation** is solved
- We construct the **dual** of this relaxation
- Duals are very similar between two nodes
- At node ν , we use this relation to hopefully **prune the child node** of ν constructed by branching on a given index



Numerical results

- **Data generation:**

- Set $(m, n) = (500, 1000)$
- Generate a random matrix $\mathbf{A} \in \mathbf{R}^{m \times n}$ with a correlation ρ between the columns
- Generate a k -sparse vector $\mathbf{x}^\dagger \in \mathbf{R}^n$
- Set $\mathbf{y} = \mathbf{A}\mathbf{x}^\dagger + \text{noise}$ with 10dB SNR
- Calibrate λ and M statistically to recover \mathbf{x}^\dagger

- **Concurrent methods:**

- Direct method using CPLEX
- Tailored BnB algorithm from Mhenni *et. al.*
- Tailored BnB algorithm with **node-screening**

	ρ	k	Direct			BnB			BnB + node-scr.		
			N	T	F	N	T	F	N	T	F
Low	5	0.1	25.9	0	0	0.1	1.5	0	0.05	0.7	0
	7	0.3	60.8	0	0	0.2	5.1	0	0.1	3.0	0
	9	0.8	102.6	10	0	0.5	15.6	0	0.4	9.8	0
High	5	1.4	10.2	0	0	1.0	6.4	0	0.7	4.2	0
	7	17.6	106.5	0	0	10.5	79.3	0	7.9	52.2	0
	9	80.7	353.4	50	0	47.8	346.4	48	41.2	267.0	40

Table: (N) : Nodes explored $\times 10^{-3}$. (T) : Solving time in sec. (F) : Number of instances not solved within 10^3 sec.

- **Observations:**

- BnB + node-scr. has the best performances
- Time gain is larger than node gain
- Double kiss-cool effect : the bounding step is performed all the faster as **many variables are fixed**

Take home message

In a BnB tailored to a sparse problem, there is not always need to perform heavy computations to prune nodes. Many nodes can be **easily pruned** by performing **simple and cheap tests** like the node-screening one.